Intro to Urban Computing - Homework 3

Christopher Williams

November 6, 2025

1   Poor Testing Model

If a linear regression model is testing poorly, there are two potential reasons that stand out immediately.

Reason 1: Overfitting. The model could be performing poorly on test data because that model has been trained *too much* on the training data, meaning it has low to no generalizability. The model has taken into account too much training data and is highly sensitive and tries to include it.

Reason 2: Underfitting. The model could also be performing poorly because the model was not trained enough on the training data, meaning that it underfit. It is not taking into account enough of the data to accurately predict, also resulting in low generalizability.

2   Linear Regression Estimation

Of the five equations listed, only 2 can be accurately estimated with linear regression. For an equation/function to be estimated accurately with linear regression, the function to be *linear in parameters*.

- **a.** $\boxed{y = \sum_{i=0}^{n} a_i x^i}$ is **estimated accurately**. This is because when it is expanded out into $y = a + a_1 x + a_2 x^2 + a_3 x^3 + ...$, it is linear in parameters and each coefficient only appears once, multiplied by a function of x.

- **b.** $\boxed{y = ax + b \cdot sin(x) + c \cdot \log(x) + d}$ is **not** estimated accurately. When expanded, terms like $sin(x)log(x)ax$ and $axcd$ appear, meaning the coefficients multiply each other which means that they are not linear, or that the coefficients cannot be separated.

- **c.** $\boxed{y = ae^{(bx)}}$ is **maybe** estimated accurately. This function could be linearized with log transformations $(\log(y) = \log(a) + bx)$
  the parameters are not linear, meaning they cannot be directly estimated with linear regression.

- **d.** $\boxed{y = ax + bx + c}$ is **estimated accurately**. Since this can simplify to $y = \beta x + c$, where $\beta = (a + b)$, this fits our general form of linear regression, the parameters are linear. We cannot separate the parameters, only find their sum, but that's still pretty good!

- **e.** $\boxed{y = \dfrac{(ax + b)}{(cx + d)}}$ is **not** estimated accurately. This does not fit in the general form and the coefficients appear in a non-linear way. And, if this function is plotted (in Desmos), unless $d = 1$, then the function cannot possibly be estimated accurately.

3   Least Squares Methods

To derive the least squares estimates for a, b, and c, I used the general form of linear regression and then used row reduction to find the least squares.

$$y = ax_1 + bx_2 + c \to y = X\beta + c \tag{1}$$

We are trying to find $X^T X$ where $X = [abc]^T$

$$X = \begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 1 \end{bmatrix} \quad y = \begin{bmatrix} 0 \\ 1.5 \\ 2.0 \\ 2.5 \end{bmatrix} \quad X^T = \begin{bmatrix} 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix} \tag{2}$$

$$X^T X = \begin{bmatrix} 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 2 & 1 & 2 \\ 1 & 2 & 2 \\ 2 & 2 & 4 \end{bmatrix} \tag{3}$$

$$X^T y = \begin{bmatrix} 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 0 \\ 1.5 \\ 2.0 \\ 2.5 \end{bmatrix} = \begin{bmatrix} 4.5 \\ 4.0 \\ 6.0 \end{bmatrix} \tag{4}$$

Now that we have our necessary equations, we need to solve $(X^T X)\beta = X^T y$

$$\begin{bmatrix} 2 & 1 & 2 \\ 1 & 2 & 2 \\ 2 & 2 & 4 \end{bmatrix} \begin{bmatrix} a \\ b \\ c \end{bmatrix} = \begin{bmatrix} 4.5 \\ 4.0 \\ 6.0 \end{bmatrix} \tag{5}$$

Original system:

$$2a + b + 2c = 4.5 \tag{1}$$
$$a + 2b + 2c = 4.0 \tag{2}$$
$$2a + 2b + 4c = 6.0 \tag{3}$$

Augmented matrix:

$$\left[\begin{array}{ccc|c} 2 & 1 & 2 & 4.5 \\ 1 & 2 & 2 & 4.0 \\ 2 & 2 & 4 & 6.0 \end{array} \right]$$

$R_1 \to \frac{1}{2}R_1$:

$$\left[\begin{array}{ccc|c} 1 & 0.5 & 1 & 2.25 \\ 1 & 2 & 2 & 4.0 \\ 2 & 2 & 4 & 6.0 \end{array} \right]$$

$R_2 \to R_2 - R_1$:

$$\left[\begin{array}{ccc|c} 1 & 0.5 & 1 & 2.25 \\ 0 & 1.5 & 1 & 1.75 \\ 2 & 2 & 4 & 6.0 \end{array} \right]$$

$R_3 \rightarrow R_3 - 2R_1$:

$$\left[\begin{array}{ccc|c} 1 & 0.5 & 1 & 2.25 \\ 0 & 1.5 & 1 & 1.75 \\ 0 & 1 & 2 & 1.5 \end{array}\right]$$

$R_2 \rightarrow \frac{1}{1.5}R_2$:

$$\left[\begin{array}{ccc|c} 1 & 0.5 & 1 & 2.25 \\ 0 & 1 & \frac{2}{3} & \frac{7}{6} \\ 0 & 1 & 2 & 1.5 \end{array}\right]$$

$R_3 \rightarrow R_3 - R_2$:

$$\left[\begin{array}{ccc|c} 1 & 0.5 & 1 & 2.25 \\ 0 & 1 & \frac{2}{3} & \frac{7}{6} \\ 0 & 0 & \frac{4}{3} & \frac{1}{3} \end{array}\right]$$

From the last row: $\frac{4}{3}c = \frac{1}{3}$, so $c = \frac{1}{4} = 0.25$

Back substitution into row 2: $b + \frac{2}{3}(0.25) = \frac{7}{6}$

$$b + \frac{1}{6} = \frac{7}{6} \implies b = 1$$

Back substitution into row 1: $a + 0.5(1) + 0.25 = 2.25$

$$a = 1.5$$

Solution: $a = 1.5$, $b = 1.0$, $c = 0.25$

4   Boston Housing Dataset

For this question, I initially output the dataframe of the dataset and looked around for any valuable predictability from the features but was having trouble with my method of finding correlation with the MEDV target through the `.corr` function.

```
# RM (avg num. rooms)
print("\nBasic Statistics:")
print(df['rm'].describe())
print(f"\nSkewness: {df['rm'].skew():.3f}")
print(f"Kurtosis: {df['rm'].kurtosis():.3f}")

correlation = df['rm'].corr(df['medv'])
print(f"\nCorrelation with MEDV (target): {correlation:.3f}")
```

which output

```
     Basic Statistics:
count    506.000000
mean       6.284634
std        0.702617
min        3.561000
25%        5.885500
50%        6.208500
75%        6.623500
max        8.780000
Name: rm, dtype: float64

```

```
12  Skewness: 0.404
13  Kurtosis: 1.892
14
15  Correlation with MEDV (target): 0.695
```

Wow, at nearly 70% correlation! That is fantastic. Lets see what that looks like in graph form.

Figure 1: Subplot of Graphs Showing Correlation of RM to MEDV

I created a graph for the distribution of RM (average number of rooms per dwelling), a box plot to find outliers, a scatterplot with the target MEDV, and a binned analysis just to see if there were any other patterns available (which there were!).

But, continuing this method with other features resulted in less exciting results.

```
1    RAD: Correlation with MEDV (target): -0.382
2    TAX: Correlation with MEDV (target): -0.469
3    INDUS: Correlation with MEDV (target): -0.484
4    CRIM: Correlation with MEDV (target): -0.388
5    ZN: Correlation with MEDV (target): 0.360
6    DIS: Correlation with MEDV (target): 0.250
```
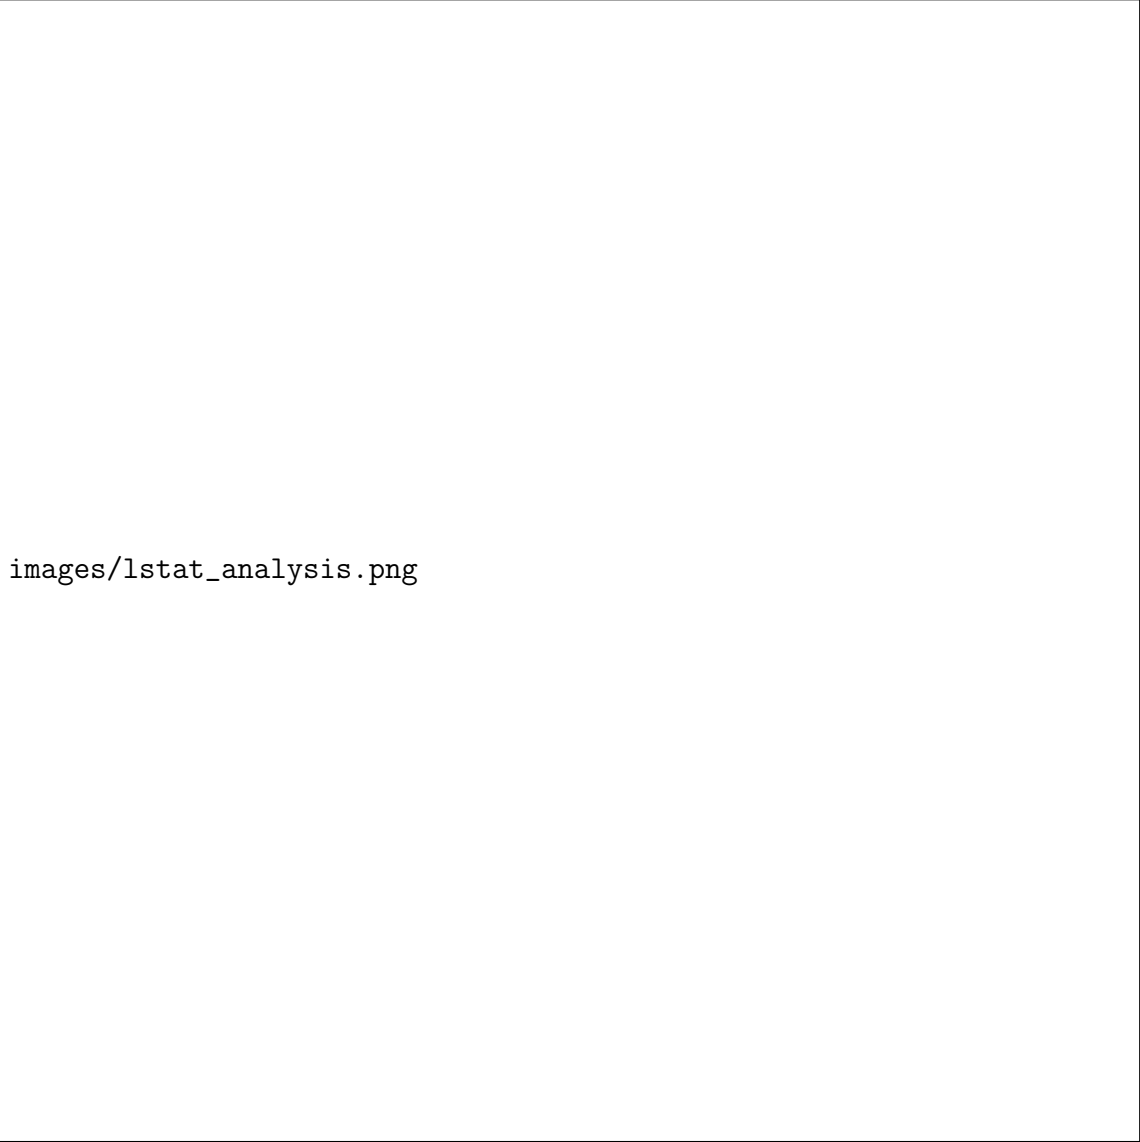
```
7     B-1000(Bk-0.63)^2: Correlation with MEDV (target): 0.333
```

Really not that exciting of findings! The best we could find in that whole debacle was the inverse correlation INDUS had with the target. Which really isn't that helpful anyways.

But then I struck gold...

```
1  LSTAT: Correlation with MEDV (target): -0.738
```

Wow! A nearly 74% inverse correlation with the target. Now we are getting somewhere. Let's visualize that relation as well.

Figure 2: Subplot of Graphs Showing Correlation of LSTAT to MEDV

However, I remembered from a machine learning course that I took that there was a method to find the best features automatically. So I did that.

4.1 Exploratory Analysis

The exploratory analysis showed that the top 5 features are LSTAT, RM, PTRATIO, TAX and INDUS. Generating a correlation matrix, feature correlation with target, target variable distribution, and

the top 3 correlated features vs. `MEDV` looks like

Figure 3: Subplot of Graphs Showing Analysis of Features

Afterwards I checked the multicolinearity between features with a high correlation and found

```
1    Multicollinearity Check (Features with |correlation| > 0.7):
2  indus <-> nox: 0.764
3  indus <-> dis: -0.708
4  indus <-> tax: 0.721
5  nox <-> age: 0.731
6  nox <-> dis: -0.769
7  age <-> dis: -0.748
8  rad <-> tax: 0.910
9
10 Train set size: 404, Test set size: 102
```

which shows that RAD and TAX, AGE and DIS, and NOX and DIS have the highest correlations among the outputs. With this understood, we can do linear regression on a subset of features chosen by the analysis. The feature selection analysis tested subsets from 1 to 13 features. Cross-validation

revealed that using all 13 features produced the highest R^2 score of 0.6688. This suggests that even weakly correlated features (like ZN with r=0.36) contribute meaningful information when combined in the linear model. The feature selection curve (Figure **??**, top-left) shows that performance plateaus after 8-10 features, with marginal gains from additional features.

4.2 Linear Regression

After applying linear regression, we get these output plots and we are mainly interested in the bottom right graph.
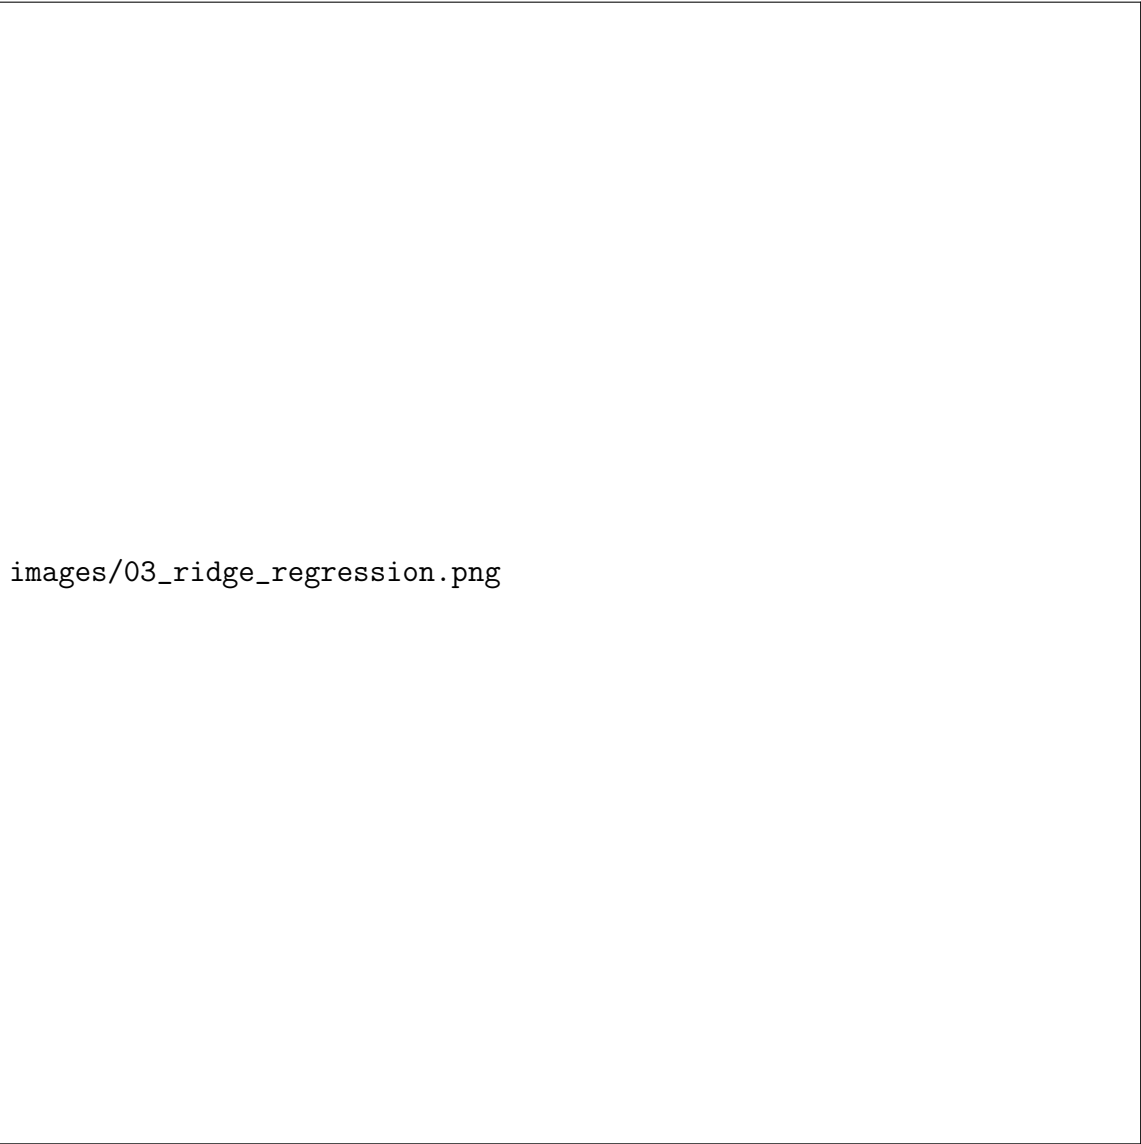

images/02_linear_regression.png

Figure 4: Subplot of Graphs Showing Linear Regression

4.3 Ridge Regression

Now, Ridge Regression (and Lasso) do not need the feature selection because they'll do it on there own! After creating a scaler object from sklearn, I found the optimal alpha via cross-validation to use, and then found the R^2 value.
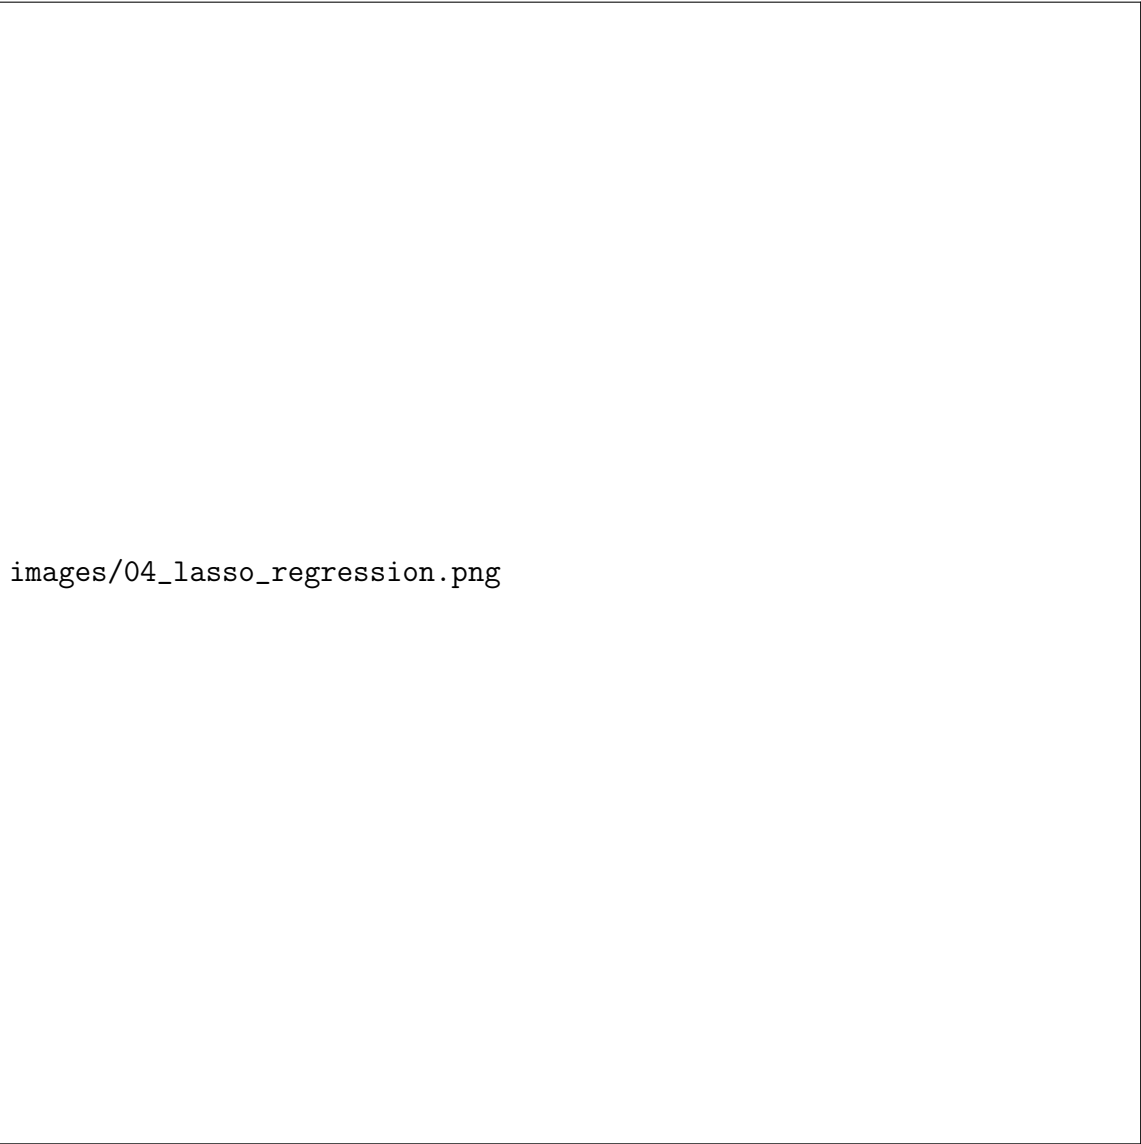
Figure 5: Subplot of Graphs Showing Ridge Regression

4.4 Lasso

Finally, we have Lasso. Again, this method doesn't need feature selection. So, running lasso after finding the optimal alpha, this is what was produced. Figure **??** shows Lasso's performance. The bottom-right subplot reveals that Lasso selected the same 9 features as Linear Regression initially identified as most important, validating our manual feature selection approach. Notably, Lasso eliminated AGE, INDUS, CHAS, and ZN, which were among the weakest predictors.

images/04_lasso_regression.png

Figure 6: Subplot of Graphs Showing Lasso

4.5   Comparison and Conclusions

Method	Features Used	Test R^2	Test RMSE	Test MAE
Linear Regression	13	0.6688	4.9286	3.1891
Ridge Regression	13	0.6679	4.9349	3.1815
Lasso Regression	13	0.6687	4.9289	3.1882

Table 1: Comparison of Regression Methods on Boston Housing Dataset

Performance Analysis:
Looking at the results, all three methods performed almost identically! The test R^2 scores ranged from 0.6679 to 0.6688, meaning each model explains about 67% of the variance in home prices. The RMSE values are basically the same too (around 4.93-4.95), which means our predictions are off by roughly $4,900 on average.

Feature Selection Results:
Here's where things got interesting - all three methods ended up using all 13 features:

- **Linear Regression**: The cross-validation kept adding features and performance kept improving, maxing out at $R^2 = 0.7244$ when using all 13. Even the weaker features like CHAS (F-score = 15.16) contributed something useful.

- **Ridge Regression**: With $\alpha = 2.848$, Ridge kept all the features but shrunk their coefficients. You can see this clearly with NOX, which went from -17.20 in Linear Regression down to -1.94 in Ridge - that's the regularization at work!

- **Lasso Regression**: The optimal $\alpha = 0.001$ was tiny, so Lasso barely regularized at all and didn't eliminate any features. Turns out that for this dataset, all the features are actually helpful and cutting any of them would hurt performance.

Handling Multicollinearity:
Remember that multicollinearity we found earlier? (RAD-TAX at 0.91, NOX-DIS at -0.77) Well, the regularization methods handled it pretty well:

- Linear Regression had some wild coefficient values (NOX = -17.20!), which is a red flag for multicollinearity issues

- Ridge smoothed everything out and gave more reasonable, stable coefficients

- Both Ridge and Lasso ended up with similar standardized coefficients, suggesting they both found good solutions to the multicollinearity problem

Overfitting Analysis:

Method	Train R^2	Test R^2	Overfit Gap
Linear Regression	0.7509	0.6688	0.0821
Ridge Regression	0.7508	0.6679	0.0829
Lasso Regression	0.7509	0.6687	0.0822

Table 2: Overfitting Analysis (Train R^2 - Test R^2)

All three models showed about the same amount of overfitting (gap of around 0.082), which is actually not bad at all. The fact that Ridge and Lasso didn't reduce overfitting much compared to regular Linear Regression suggests our dataset is big enough (506 samples) for the 13 features we're using.

Most Important Features:
Looking at the coefficient magnitudes across all methods, the top 5 features that matter most are:

1. **LSTAT** (lower status %): The strongest predictor at -3.61

2. **RM** (rooms): Close second at +3.15

3. **DIS** (distance to employment): -3.07

4. **RAD** (highway access): +2.23

5. **PTRATIO** (pupil-teacher ratio): -2.04

This matches perfectly with what we found in the exploratory analysis - LSTAT (-0.738 correlation) and RM (+0.695 correlation) were our stars from the beginning!

Key Takeaways:

- Even with all that multicollinearity, every feature ended up being useful

- Lasso's tiny alpha (0.001) tells us that cutting features would actually make things worse for this dataset

- Ridge's coefficient shrinkage really helped stabilize things - just look at how NOX went from -17.20 to -1.94!

- The nearly identical performance across all three methods means this dataset plays nice with linear models and doesn't need heavy regularization

5   Code

I think I got it working this time. Thank you TA's for your hard work. Please enjoy your week(end) :)

<div align="center">

THE CODE

</div>